

Batch 1: Questions 1–20

Questions

Question 1

An organization is developing an enterprise AI governance program. The leadership team wants to distinguish strategic oversight from day-to-day execution. Which activity BEST represents **AI governance** rather than AI management?

- A. Configuring production infrastructure for an AI application
 - B. Assigning developers to implement a model-monitoring pipeline
 - C. Establishing organizational AI risk tolerance and oversight mechanisms
 - D. Troubleshooting a failed model deployment
-

Question 2

A multinational organization operates AI systems in the EU, United States, and several Asian markets. Its leadership wants one overarching governance approach that minimizes the need for completely different control frameworks in every jurisdiction. Which approach BEST aligns with the Module 6 guidance?

- A. Apply only the regulations of the organization's headquarters
 - B. Select the least restrictive jurisdiction and use it globally
 - C. Comply with the highest applicable standards while maintaining operational efficiency
 - D. Develop completely independent governance programs for every country
-

Question 3

Which statement BEST describes the fundamental difference between traditional technology governance and AI governance?

- A. Traditional technology governance focuses on data, while AI governance focuses only on algorithms
 - B. AI governance eliminates the need for cybersecurity controls
 - C. AI governance additionally addresses risks inherent in AI behavior, including fairness, predictability, safety, and societal impact
 - D. Traditional technology governance is mandatory, whereas AI governance is always voluntary
-

Question 4

A U.S.-based company provides an AI-powered service to customers located in the European Union. The company's executives argue that the EU AI Act cannot apply because the organization has no physical office in Europe. Which response BEST reflects the Module 6 discussion?

- A. The executives are correct because AI regulations only apply where the provider is incorporated
- B. The EU AI Act can have extraterritorial reach when AI systems affect EU residents or markets
- C. The EU AI Act applies only to organizations headquartered in EU member states
- D. The EU AI Act applies only when the AI system processes financial information

Question 5

Which regulatory approach BEST characterizes the United Kingdom's AI governance model described in Module 6?

- A. A single comprehensive AI regulator with centralized authority
 - B. A principles-based approach relying on existing sector regulators
 - C. Complete technology-industry self-regulation without government involvement
 - D. A comprehensive legislative framework identical to the EU AI Act
-

Question 6

A financial institution wants to understand which organization would have sector-specific responsibility for AI-related oversight in UK financial services under the model described in Module 6. Which organization is MOST relevant?

- A. Ofcom
 - B. Competition and Markets Authority
 - C. Financial Conduct Authority
 - D. European AI Office
-

Question 7

An organization is implementing the NIST AI RMF and begins by identifying the AI system's context, categorizing its risks, and analyzing potential stakeholder impacts. Which NIST AI RMF function is being implemented?

- A. GOVERN
 - B. MAP
 - C. MEASURE
 - D. MANAGE
-

Question 8

An AI model monitoring team continuously evaluates model drift, bias, performance degradation, and security vulnerabilities. Which NIST AI RMF function MOST directly encompasses these activities?

- A. GOVERN
 - B. MAP
 - C. MEASURE
 - D. MANAGE
-

Question 9

A business determines that an AI risk should be mitigated rather than accepted. The organization allocates resources, implements remediation procedures, and establishes an incident response process. Which NIST AI RMF function BEST describes these activities?

- A. GOVERN
 - B. MAP
 - C. MEASURE
 - D. MANAGE
-

Question 10

An AI system consistently produces accurate results during development but begins producing poorer results after several months because the characteristics of incoming data have changed. Which concept from Module 6 BEST describes this situation?

- A. Data lineage
 - B. Model drift
 - C. Regulatory arbitrage
 - D. Conformity assessment
-

Question 11

According to the NIST AI RMF trustworthiness characteristics discussed in Module 6, which characteristic specifically focuses on preventing harm to individuals, organizations, or society from intended and unintended consequences?

- A. Fair
 - B. Safe
 - C. Transparent
 - D. Privacy-enhanced
-

Question 12

A bank wants an internationally recognized AI risk-management methodology that can integrate with existing ISO-based management systems. It does not specifically require direct certification from the risk-management standard itself. Which framework is the BEST fit?

- A. MITRE ATLAS
 - B. ISO/IEC 23894
 - C. NIST AI RMF
 - D. CSA AI-CAIQ
-

Question 13

A security operations team wants a framework specifically designed to help identify, analyze, and defend against adversarial threats targeting AI systems. Which framework should the organization prioritize?

- A. ISO/IEC 23894
 - B. NIST AI RMF
 - C. MITRE ATLAS
 - D. ISO 42001
-

Question 14

An organization wants a **certifiable AI management system** aligned with internationally recognized requirements. Which framework is MOST appropriate according to the Module 6 framework-selection matrix?

- A. NIST AI RMF
 - B. ISO/IEC 23894
 - C. ISO 42001
 - D. MITRE ATLAS
-

Question 15

A cloud service provider wants to assess its AI security controls and eventually demonstrate assurance through the Cloud Security Alliance's STAR for AI program. Which sequence BEST reflects the integrated CSA approach described in Module 6?

- A. STAR certification → AI-CAIQ → AICM controls
 - B. AI-CAIQ → AICM controls → STAR assurance/certification
 - C. MITRE ATLAS → ISO 42001 → AI-CAIQ
 - D. NIST AI RMF → GDPR → STAR certification
-

Question 16

An organization is establishing an AI Safety Officer position. AI is considered a core strategic capability and represents significant enterprise risk. Which reporting structure is identified in Module 6 as providing the **maximum organizational authority and resource access**?

- A. Dual-hat CTO/CDO role
 - B. Chief Information Security Officer report
 - C. Chief Risk Officer report
 - D. Direct CEO/COO report
-

Question 17

An organization is designing a RACI matrix for AI security controls. Which group would MOST appropriately be assigned **Responsible** for implementation activities?

- A. Executive leadership and the board
- B. AI security engineers, DevSecOps teams, and operations analysts
- C. External regulators and auditors
- D. The Chief Risk Officer exclusively

Question 18

An AI governance committee is reviewing its RACI matrix. Several activities currently have multiple individuals marked as "Accountable." What is the MOST important governance issue the committee should address?

- A. Increase the number of people marked Responsible
 - B. Ensure each activity has a single point of accountability
 - C. Eliminate all Consulted stakeholders
 - D. Move all accountability to technical teams
-

Question 19

An organization is preparing a risk card for a high-impact AI system. Which information belongs MOST directly in the **Human Involvement and Control Mechanisms** section?

- A. Model version, modality, and technical architecture
 - B. Accuracy metrics across different populations
 - C. Human oversight requirements, end-user control capabilities, and feedback mechanisms
 - D. Likelihood and severity ratings for fundamental-rights impacts
-

Question 20

A multinational organization wants to implement AI governance using several frameworks simultaneously. Its compliance team proposes selecting whichever framework is easiest to implement in each country, even when another applicable framework has stricter requirements. Which approach BEST aligns with Module 6?

- A. Always use the least restrictive framework to reduce operational complexity
 - B. Select frameworks solely according to implementation cost
 - C. Comply with the highest applicable standards while maintaining operational efficiency
 - D. Use only one framework globally regardless of regulatory requirements
-

Answers & Explanations

The questions above intentionally draw from different parts of Module 6 rather than repeatedly testing the same concepts. The study guide explicitly frames Module 6 around governance, regulatory requirements, risk frameworks, organizational accountability, documentation, monitoring, cross-border issues, and implementation.

Question 1

Correct Answer: C — Establishing organizational AI risk tolerance and oversight mechanisms

Explanation:

AI governance provides **strategic oversight and direction**, including risk tolerance, policy development, accountability, performance evaluation, and compliance oversight. AI management is

more concerned with operational execution, implementation, resource allocation, and day-to-day administration.

Question 2

Correct Answer: C — Comply with the highest applicable standards while maintaining operational efficiency

Explanation:

Module 6 recommends that organizations operating across jurisdictions develop governance frameworks capable of complying with the **highest applicable standards**, while avoiding unnecessary operational complexity. This is particularly important because AI regulations vary significantly across jurisdictions.

Question 3

Correct Answer: C — AI governance additionally addresses risks inherent in AI behavior, including fairness, predictability, safety, and societal impact

Explanation:

Traditional technology governance generally focuses on protecting systems and data from internal and external threats. AI governance expands the scope to include risks inherent in the AI system itself, such as unpredictable behavior, fairness, safety, societal impact, and regulatory compliance.

Question 4

Correct Answer: B — The EU AI Act can have extraterritorial reach when AI systems affect EU residents or markets

Explanation:

The Module 6 study guide explicitly describes the EU AI Act as having **extraterritorial reach**. Organizations outside the EU can therefore face obligations when their AI systems affect EU residents or markets.

Question 5

Correct Answer: B — A principles-based approach relying on existing sector regulators

Explanation:

The UK approach described in Module 6 is **principles-based and distributed**. Rather than establishing one centralized AI regulator or adopting a comprehensive AI-specific statute equivalent to the EU AI Act, existing regulators adapt their oversight capabilities to AI-specific risks.

Question 6

Correct Answer: C — Financial Conduct Authority

Explanation:

Module 6 identifies the **Financial Conduct Authority (FCA)** as the relevant existing regulator for AI-related oversight within UK financial services. Ofcom is associated with telecommunications, while the Competition and Markets Authority addresses competition-related impacts.

Question 7

Correct Answer: B — MAP

Explanation:

The **MAP** function identifies AI-system context, categorizes risks, and characterizes potential impacts. It includes activities such as AI system inventory, risk classification, and stakeholder impact analysis.

Question 8

Correct Answer: C — MEASURE

Explanation:

The **MEASURE** function analyzes, assesses, and monitors AI risks. The study guide specifically identifies model drift, bias detection, performance degradation, and security vulnerabilities as examples of activities within this function.

Question 9

Correct Answer: D — MANAGE

Explanation:

The **MANAGE** function implements risk-response strategies. These include risk acceptance, mitigation, transfer, avoidance, incident response, remediation, and continuous improvement.

Question 10

Correct Answer: B — Model drift

Explanation:

Model drift is the decline in model performance over time caused by changes in data patterns. Module 6 emphasizes continuous monitoring because AI systems can degrade even when they initially performed adequately.

Question 11

Correct Answer: B — Safe

Explanation:

The NIST AI RMF **Safe** characteristic addresses avoiding harm to individuals, organizations, or society

resulting from both intended and unintended consequences. The study guide also connects safety assessment with failure-mode analysis, impact analysis, and safeguards.

Question 12

Correct Answer: B — ISO/IEC 23894

Explanation:

ISO/IEC 23894 provides internationally recognized guidance specifically for **AI risk management** throughout the AI life cycle. It is based on ISO 31000 principles and is designed to integrate with existing ISO management systems.

A key exam distinction is:

- **ISO 23894** → AI risk-management guidance
- **ISO 42001** → certifiable AI management-system requirements

The module explicitly makes this distinction.

Question 13

Correct Answer: C — MITRE ATLAS

Explanation:

MITRE ATLAS extends the MITRE ATT&CK methodology to AI systems and provides a structured approach to AI-specific threat intelligence, red teaming, penetration testing, threat hunting, and incident response.

Question 14

Correct Answer: C — ISO 42001

Explanation:

The Module 6 framework-selection matrix identifies **ISO 42001** as an AI management-system framework with certification available. ISO 23894, by contrast, provides risk-management guidance and does not itself provide direct certification.

Question 15

Correct Answer: B — AI-CAIQ → AICM controls → STAR assurance/certification

Explanation:

The CSA integrated approach described in Module 6 follows a logical progression:

1. Use **AI-CAIQ** to establish a governance baseline.
2. Implement **AICM controls**.
3. Pursue **STAR** assurance/certification to demonstrate trustworthiness.

The study guide explicitly presents this sequence.

Question 16

Correct Answer: D — Direct CEO/COO report

Explanation:

The Module 6 organizational structure identifies a **Direct CEO/COO Report** as appropriate where AI is a core strategic capability or represents significant enterprise risk. The advantage is maximum organizational authority and resource access.

Question 17

Correct Answer: B — AI security engineers, DevSecOps teams, and operations analysts

Explanation:

The **Responsible** role in a RACI matrix performs or implements the activity. The sample questions you supplied specifically identify AI security engineers, DevSecOps teams, and operations analysts as Responsible for AI security controls.

Question 18

Correct Answer: B — Ensure each activity has a single point of accountability

Explanation:

An important RACI governance principle is maintaining **one clear Accountable party** for an activity. Multiple people can participate in implementation or consultation, but accountability should not become ambiguous. The sample questions explicitly emphasize a single "A" per activity and balanced workload distribution.

Question 19

Correct Answer: C — Human oversight requirements, end-user control capabilities, and feedback mechanisms

Explanation:

The **Human Involvement and Control Mechanisms** section of a risk card includes end-user interaction patterns, control capabilities, human oversight requirements, subject involvement/consent, feedback mechanisms, and complaint handling.

Question 20

Correct Answer: C — Comply with the highest applicable standards while maintaining operational efficiency

Explanation:

Module 6 explicitly recommends a multi-framework approach that meets the **highest applicable**

requirements while avoiding unnecessary operational complexity. This is particularly important for multinational organizations subject to overlapping regulatory and framework requirements.

Questions

Question 21

A U.S. organization is developing an AI governance program and asks whether the NIST AI RMF creates mandatory federal legal obligations for every organization using AI. Which response BEST reflects the Module 6 study guide?

- A. Yes, NIST AI RMF is a mandatory federal regulation for all AI systems
 - B. Yes, but only organizations processing personal information must comply
 - C. No, NIST AI RMF is voluntary guidance, although it can provide a structured approach to AI risk management
 - D. No, because NIST AI RMF applies only to government agencies
-

Question 22

A company markets an AI-powered facial recognition product and claims that it has substantially higher accuracy than competing products. The company has insufficient evidence to substantiate the claim. Which U.S. regulatory body and enforcement principle are MOST relevant according to Module 6?

- A. FDA — medical device conformity assessment
 - B. FTC — accountability for deceptive AI claims
 - C. DHS — export-control enforcement
 - D. NIST — mandatory AI certification
-

Question 23

Under the California CCPA Automated Decision-Making Technology (ADMT) requirements discussed in Module 6, when should an organization provide the required notice to consumers?

- A. Only after the AI system has made an automated decision
 - B. Before processing personal information through the ADMT
 - C. Only when the consumer submits a complaint
 - D. Only when the AI system processes sensitive personal information
-

Question 24

A California consumer is subject to an automated decision made by an AI system and wants meaningful information about how that system affects the outcome. Which combination of rights is MOST consistent with the Module 6 discussion of ADMT?

- A. Right to opt out, right to access information, right to human review, and right to explanation
 - B. Right to delete the AI model and right to source-code ownership
 - C. Right to prohibit all automated processing globally
 - D. Right to require the organization to disclose proprietary model weights
-

Question 25

A company argues that consumer-protection laws should not apply to its AI product because "AI is a new technology and therefore existing consumer-protection laws do not cover it." Which statement BEST reflects the FTC position described in Module 6?

- A. AI systems are exempt from existing consumer-protection laws
 - B. Only open-source AI is subject to consumer-protection laws
 - C. There is no AI exemption from existing consumer-protection laws
 - D. Consumer-protection laws apply only after an AI-specific federal law is enacted
-

Question 26

Which of the following BEST describes the U.S. AI regulatory environment presented in Module 6?

- A. A single comprehensive federal AI law administered by one regulator
 - B. A distributed framework involving federal policy, agency enforcement, and state-level requirements
 - C. A voluntary industry-only framework with no government involvement
 - D. A regulatory structure identical to the EU AI Act
-

Question 27

An organization wants to monitor U.S. federal AI policy developments because executive orders and associated implementation guidance may change as policy priorities evolve. What does Module 6 recommend?

- A. Ignore federal developments until Congress passes a comprehensive AI statute
 - B. Rely exclusively on industry blogs for regulatory updates
 - C. Monitor official government channels and consult qualified legal counsel
 - D. Apply the EU AI Act instead of monitoring U.S. developments
-

Question 28

Under the EU AI Act implementation timeline presented in Module 6, which requirement became effective beginning **February 2, 2025**?

- A. Comprehensive high-risk AI system requirements
- B. General-purpose AI model requirements
- C. Prohibitions on specified unacceptable-risk AI systems
- D. Extended Annex I high-risk requirements

Question 29

An organization operating an AI system in the EU discovers that its system uses subliminal techniques capable of causing psychological harm. What should the organization do under the Module 6 description of the EU AI Act?

- A. Continue operating the system if human oversight is added
 - B. Continue operating it after completing a risk assessment
 - C. Immediately discontinue the prohibited use
 - D. Wait until August 2026 when high-risk requirements take effect
-

Question 30

Which AI application is identified in Module 6 as an example of a **prohibited AI system** under the EU AI Act?

- A. An AI system recommending products to consumers
 - B. A biometric categorization system that infers sensitive attributes such as political opinions
 - C. An internal employee scheduling system
 - D. An AI system summarizing publicly available documents
-

Question 31

A company deploys an AI system that performs real-time remote biometric identification in public spaces. Under the Module 6 description of prohibited AI systems, which statement is **MOST** accurate?

- A. Such systems are always permitted if the organization publishes an AI transparency notice
 - B. Such systems are prohibited, subject to limited law-enforcement exceptions
 - C. Such systems are automatically classified as low-risk
 - D. Such systems are exempt from the EU AI Act if developed outside Europe
-

Question 32

A company is preparing for the EU AI Act requirements that apply to high-risk AI systems. Which activity is **MOST** directly associated with the high-risk requirements described in Module 6?

- A. Eliminating all human involvement from AI decision-making
 - B. Conducting pre-deployment conformity assessment and maintaining technical documentation
 - C. Using only open-source AI models
 - D. Avoiding post-market monitoring once the system is deployed
-

Question 33

A high-risk AI provider makes a significant improvement to its model's capabilities after completing its initial compliance documentation. What does Module 6 indicate about the provider's documentation obligations?

- A. Documentation is required only once, before the first deployment
 - B. Documentation must be maintained throughout the lifecycle and updated to reflect significant changes
 - C. Documentation is necessary only if a regulator requests it
 - D. Only the model's source code needs to be updated
-

Question 34

Which activity BEST represents the **post-market surveillance** obligation described for high-risk AI systems under Module 6?

- A. Monitoring system performance after deployment and reporting relevant incidents or safety concerns
 - B. Performing model development only before initial deployment
 - C. Eliminating all production logging to protect user privacy
 - D. Replacing the AI system every six months regardless of performance
-

Question 35

A foundation model exceeds the computational threshold identified in Module 6 for systemic-risk models. Which additional activity is specifically associated with such models?

- A. Eliminating all downstream use of the model
 - B. Conducting systemic-risk assessments covering misuse, security vulnerabilities, and societal impact
 - C. Obtaining automatic exemption from EU AI requirements
 - D. Treating the model as a low-risk application because it is general-purpose
-

Question 36

According to Module 6, what training-computation threshold is associated with the systemic-risk general-purpose AI models discussed in the study guide?

- A. 10^{23} FLOPs
 - B. 10^{24} FLOPs
 - C. 10^{25} FLOPs
 - D. 10^{26} FLOPs
-

Question 37

An organization violates the EU AI Act's provisions concerning a prohibited AI system. According to the penalty structure presented in Module 6, which penalty level is associated with such violations?

- A. Up to €7.5 million or 1.5% of global annual turnover, whichever is higher
 - B. Up to €15 million or 3% of global annual turnover, whichever is higher
 - C. Up to €35 million or 7% of global annual turnover, whichever is higher
 - D. A mandatory fixed penalty of €1 million
-

Question 38

A cloud-based AI provider wants to use the CSA AICM as a structured mechanism for evaluating its AI security posture. Which description BEST matches the AICM in Module 6?

- A. A threat-intelligence framework containing adversarial attack techniques
 - B. A cloud and AI security-control framework designed for AI systems
 - C. A certifiable AI management-system standard
 - D. A privacy regulation governing consumer data
-

Question 39

According to Module 6, how is **STAR for AI Level 1** primarily performed?

- A. An independent regulator conducts a mandatory legal investigation
 - B. Organizations use the AI-CAIQ to self-attest their implementation of AICM controls
 - C. An EU-notified body performs a conformity assessment
 - D. An external auditor performs a SOC-based audit
-

Question 40

An organization wants independent third-party assurance against AICM-based criteria under the CSA STAR for AI model. Which level is MOST appropriate?

- A. STAR for AI Level 0
 - B. STAR for AI Level 1
 - C. STAR for AI Level 2
 - D. AI-CAIQ Level 3
-

Answers & Explanations

Question 21

Correct Answer: C — No, NIST AI RMF is voluntary guidance, although it can provide a structured approach to AI risk management

Explanation:

The Module 6 study guide describes the NIST AI RMF as **voluntary guidance**. It provides a structured framework for managing AI risk but is not itself a mandatory federal regulation applicable to every organization.

TAISE exam point:

Do not automatically equate a framework with a regulation. A framework may provide the **control methodology**, while legislation or regulation establishes the **legal obligation**.

Question 22

Correct Answer: B — FTC — accountability for deceptive AI claims

Explanation:

The FTC's **Operation AI Comply** focuses on problematic AI practices, including false or unsupported AI claims. Module 6 specifically discusses companies that made unsupported claims regarding facial recognition, security screening, and AI legal capabilities.

Question 23

Correct Answer: B — Before processing personal information through the ADMT

Explanation:

The California ADMT requirements described in Module 6 include a **pre-use notice obligation**. The notice should explain the AI system's purpose and functionality, disclose relevant consumer rights, and be provided before processing personal information.

Question 24

Correct Answer: A — Right to opt out, right to access information, right to human review, and right to explanation

Explanation:

Module 6 specifically identifies four important consumer rights in the ADMT framework:

- Right to opt out
 - Right to access information about automated decision-making
 - Right to human review
 - Right to an explanation of how AI affects consumer outcomes
-

Question 25

Correct Answer: C — There is no AI exemption from existing consumer-protection laws

Explanation:

The study guide explicitly states that the FTC's position is that there is **no AI exemption from existing consumer-protection laws**. The fact that a product uses AI does not remove existing legal obligations.

Question 26

Correct Answer: B — A distributed framework involving federal policy, agency enforcement, and state-level requirements

Explanation:

Module 6 contrasts the U.S. and EU approaches. The U.S. landscape is described as **distributed**, relying primarily on executive guidance, agency enforcement of existing laws, and state-level legislation.

Question 27

Correct Answer: C — Monitor official government channels and consult qualified legal counsel

Explanation:

The study guide warns that U.S. federal AI policy can evolve as administrations and policy priorities change. Organizations should therefore monitor official government sources and obtain appropriate legal advice rather than assuming that today's policy environment will remain unchanged.

Question 28

Correct Answer: C — Prohibitions on specified unacceptable-risk AI systems

Explanation:

The Module 6 timeline identifies **February 2, 2025** as the first major EU AI Act implementation phase. It introduced immediate restrictions on prohibited AI systems.

The timeline presented in the study guide is:

- **February 2025** → prohibited AI systems
 - **August 2025** → GPAI requirements
 - **August 2026** → most high-risk requirements
 - **August 2027** → extended timeline for specified Annex I sectors
-

Question 29

Correct Answer: C — Immediately discontinue the prohibited use

Explanation:

The study guide describes prohibited AI systems as subject to **immediate compliance requirements** beginning February 2, 2025. Organizations should cease prohibited uses rather than attempt to mitigate them through ordinary risk controls.

Important exam distinction:

For prohibited AI, the question is generally **not** "How do we mitigate the risk?" but rather "Is this use allowed at all?"

Question 30

Correct Answer: B — A biometric categorization system that infers sensitive attributes such as political opinions

Explanation:

Module 6 explicitly identifies biometric categorization systems that infer sensitive attributes—including political opinions, religious beliefs, race, and sexual orientation—as prohibited uses.

Question 31

Correct Answer: B — Such systems are prohibited, subject to limited law-enforcement exceptions

Explanation:

The study guide lists **real-time remote biometric identification in public spaces** among prohibited AI systems while recognizing limited law-enforcement exceptions.

Question 32

Correct Answer: B — Conducting pre-deployment conformity assessment and maintaining technical documentation

Explanation:

High-risk AI systems face extensive lifecycle requirements. Module 6 specifically identifies:

- Pre-deployment conformity assessment
 - Technical documentation
 - Risk management
 - Data governance
 - Transparency
 - Human oversight
 - Accuracy and robustness
 - Security measures
-

Question 33

Correct Answer: B — Documentation must be maintained throughout the lifecycle and updated to reflect significant changes

Explanation:

High-risk AI documentation is **not a one-time compliance exercise**. Module 6 states that technical documentation must be maintained throughout the AI system lifecycle and updated when significant changes or improvements affect system capabilities.

Question 34

Correct Answer: A — Monitoring system performance after deployment and reporting relevant incidents or safety concerns

Explanation:

Post-market surveillance requires ongoing monitoring after deployment. The Module 6 study guide connects this obligation with monitoring system performance and reporting relevant incidents or safety concerns to authorities.

Question 35

Correct Answer: B — Conducting systemic-risk assessments covering misuse, security vulnerabilities, and societal impact

Explanation:

Module 6 identifies additional requirements for general-purpose models with systemic-risk potential. These include comprehensive systemic-risk assessments addressing **misuse, security vulnerabilities, and societal impacts** across deployment scenarios.

Question 36

Correct Answer: C — 10^{25} FLOPs

Explanation:

The Module 6 study guide identifies **10^{25} FLOPs of training computation** as the threshold associated with the systemic-risk models discussed in this section.

Exam tip:

This is the type of numerical detail worth memorizing because it can appear as a direct knowledge question or be embedded in a scenario.

Question 37

Correct Answer: C — Up to €35 million or 7% of global annual turnover, whichever is higher

Explanation:

The penalty structure presented in Module 6 distinguishes among categories of violations:

Violation	Maximum penalty described
Prohibited AI systems	€35M or 7%
High-risk violations	€15M or 3%
Documentation/procedural violations	€7.5M or 1.5%

The prohibited-AI category therefore carries the highest penalty described in the study guide.

Question 38

Correct Answer: B — A cloud and AI security-control framework designed for AI systems

Explanation:

The CSA AICM is designed specifically to address **cloud and AI security controls**. Module 6 distinguishes it from MITRE ATLAS, which focuses on the AI threat landscape, and ISO 42001, which provides certifiable AI management-system requirements.

The study guide further describes AICM as containing **243 control objectives across 18 security domains**.

Question 39

Correct Answer: B — Organizations use the AI-CAIQ to self-attest their implementation of AICM controls

Explanation:

STAR for AI Level 1 uses the **AI-CAIQ** for organizational self-attestation of AICM implementation. It is therefore an organizational self-assessment rather than independent third-party certification.

Question 40

Correct Answer: C — STAR for AI Level 2

Explanation:

STAR for AI Level 2 introduces independent third-party assessment against AICM-based criteria, using either ISO-based or SOC-based audit schemes.

A useful progression to remember is:

AI-CAIQ → AICM implementation → STAR assurance

The Module 6 study guide explicitly presents this as the CSA pathway toward demonstrating AI trustworthiness.

Question 41

An organization is establishing an AI Governance Committee to oversee enterprise AI risks. Which committee composition BEST aligns with the Module 6 guidance?

- A. Two members exclusively from the IT department
 - B. Seven to nine core members representing varied organizational perspectives
 - C. Twenty-five members from every business function
 - D. One executive who makes all AI governance decisions independently
-

Question 42

A proposed AI system has potentially significant impacts on customers and fundamental rights. The business team wants the project approved immediately because it is commercially important. Which governance principle BEST supports requiring additional oversight before approval?

- A. AI governance should prioritize business value over risk classification
 - B. Governance decisions should be proportionate to the AI system's risk classification
 - C. All AI systems should receive identical approval procedures
 - D. High-impact AI systems should bypass governance to accelerate innovation
-

Question 43

An AI Governance Committee is reviewing a system classified as low risk. Which decision-making approach is MOST consistent with the tiered governance model described in Module 6?

- A. Require board-level approval for every low-risk operational change
 - B. Apply proportionate governance and delegate lower-risk decisions to appropriate operational levels
 - C. Eliminate all documentation requirements for low-risk systems
 - D. Require external regulatory approval before deployment
-

Question 44

An AI system has the potential to cause public harm, trigger regulatory enforcement, and create substantial reputational damage. Under the Module 6 escalation model, which escalation level is MOST appropriate?

- A. Level 1 — Operational Issues
 - B. Level 2 — Governance Violations
 - C. Level 3 — Strategic and Crisis Issues
 - D. No escalation unless the model completely fails
-

Question 45

An AI operations team detects minor model-performance degradation during routine monitoring. There is no evidence of customer harm, regulatory impact, or major security concerns. According to the Module 6 escalation model, what is the MOST appropriate initial response?

- A. Immediately escalate to the board
 - B. Treat it as a Level 1 operational issue
 - C. Initiate external regulatory notification automatically
 - D. Suspend the entire enterprise AI program
-

Question 46

An organization identifies a policy violation involving an AI vendor and a potential compliance issue. Which escalation level and responsible stakeholders BEST match the Module 6 model?

- A. Level 1; system administrators only
- B. Level 2; AI governance committee, compliance officers, and legal counsel
- C. Level 3; board committee only
- D. Level 1; external auditors only

Question 47

A Level 1 AI operational issue continues despite technical remediation and begins generating customer complaints. What does the Module 6 escalation model indicate?

- A. The issue should remain permanently at Level 1
- B. Persistent issues and customer complaints can trigger escalation
- C. The issue should automatically be classified as a prohibited AI system
- D. The issue should be ignored unless regulators contact the organization

Question 48

A serious AI incident attracts media attention and creates potential legal liability. Which authority is associated with Level 3 escalation in the Module 6 framework?

- A. Technical parameter adjustment only
- B. Routine maintenance and service configuration
- C. Strategic decisions, external communications, legal action, and business continuity
- D. Vendor onboarding and procurement approval only

Question 49

An organization is developing a model card for an AI model. Which information would BEST belong in the **Model Overview and Intended Use** section?

- A. Only the model's cybersecurity vulnerabilities
- B. Model purpose, intended applications, target users, and deployment context
- C. Only the organization's annual AI governance budget
- D. Only regulatory penalties applicable to the organization

Question 50

A model card is intended to help stakeholders understand an AI system before using or evaluating it. Which statement BEST describes its governance value?

- A. It replaces the need for risk assessments
- B. It provides structured information about the model, its intended use, limitations, and relevant characteristics
- C. It serves exclusively as a marketing document
- D. It eliminates the need for ongoing model monitoring

Question 51

An organization is documenting a dataset used to train a high-impact AI model. Which information is MOST important for identifying potential representation and quality issues?

- A. Only the name of the database administrator
 - B. Data instance types, representativeness, missing information, demographic coverage, and geographic coverage
 - C. Only the model's inference latency
 - D. Only the vendor's annual revenue
-

Question 52

An AI development team wants future users of a dataset to understand how the data was collected and what biases may have been introduced during collection. Which documentation area should they prioritize?

- A. Collection Process Documentation
 - B. Model Deployment Dashboard
 - C. Incident Escalation Register
 - D. Financial Performance Report
-

Question 53

A dataset has undergone normalization, transformation, error correction, and other preprocessing steps. Why should these procedures be documented?

- A. To prevent any future use of the dataset
 - B. To support reproducibility and help downstream users understand potential effects on model behavior and performance
 - C. To replace all data-quality testing
 - D. To ensure that the dataset can never be modified again
-

Question 54

An organization is documenting a dataset containing sensitive information. Which combination BEST represents appropriate **Distribution and Access Controls** documentation?

- A. Access procedures, licensing restrictions, intellectual-property considerations, and retention/deletion policies
 - B. Only the model's accuracy and latency
 - C. Only the names of developers who created the model
 - D. Only the AI system's regulatory classification
-

Question 55

An organization is completing an AI risk card for a system used to make decisions affecting individuals. Which information BEST belongs in the **Risk Profile Assessment** section?

- A. Only the AI system's software version
- B. Likelihood and severity, health and safety risks, fundamental-rights impacts, societal/environmental effects, and residual risk

- C. Only the names of the AI developers
 - D. Only the system's cloud hosting provider
-

Question 56

An organization wants to demonstrate that humans retain meaningful control over an AI-assisted decision-making process. Which risk-card information would BEST support this objective?

- A. Model version and technical modality only
 - B. End-user interaction, human oversight requirements, consent procedures, feedback mechanisms, and complaint handling
 - C. Only model training compute
 - D. Only the organization's cybersecurity policy
-

Question 57

An organization is beginning development of an enterprise AI policy. Which activity should occur FIRST according to the Module 6 policy development methodology?

- A. Publish the policy immediately and collect feedback afterward
 - B. Conduct comprehensive stakeholder engagement
 - C. Select a specific AI vendor
 - D. Develop only technical security requirements
-

Question 58

A cross-functional AI policy working group is being established. Which group of stakeholders BEST reflects the Module 6 recommendation?

- A. Only software developers and data scientists
 - B. Legal, compliance, IT, relevant business units, HR, and potentially affected external stakeholders
 - C. Only executive leadership
 - D. Only external auditors
-

Question 59

An organization wants its AI policy to remain useful as AI technologies evolve and new use cases emerge. Which policy-development approach BEST aligns with Module 6?

- A. Create highly rigid technology-specific requirements that cannot change
 - B. Use principle-based policies combined with specific requirements for critical risk and compliance areas
 - C. Avoid documenting requirements to preserve flexibility
 - D. Delegate policy development entirely to individual developers
-

Question 60

An AI governance team is designing a comprehensive audit trail. Which combination BEST represents the types of activities that should be captured?

- A. Only successful model predictions
 - B. Policy and configuration changes, operational activity, incidents, exceptions, remediation, and regulatory inquiries
 - C. Only employee login records
 - D. Only financial transactions generated by AI systems
-

Answers & Explanations

Question 41

Correct Answer: B — Seven to nine core members representing varied organizational perspectives

Explanation:

Module 6 identifies **seven to nine core members from varied backgrounds** as the optimal composition for an AI governance committee. The purpose is to bring together the different perspectives necessary for AI governance rather than treating AI governance as solely a technology function.

Question 42

Correct Answer: B — Governance decisions should be proportionate to the AI system's risk classification

Explanation:

Module 6 emphasizes **risk classification and proportionate oversight**. Higher-risk systems require stronger governance measures, while lower-risk systems can generally be handled through more streamlined decision processes. This supports efficient resource allocation without weakening governance over significant risks.

Question 43

Correct Answer: B — Apply proportionate governance and delegate lower-risk decisions to appropriate operational levels

Explanation:

The module teaches a **tiered decision-making framework based on risk classification**. The objective is not to send every AI decision to the highest governance level, but to align decision authority with risk.

Question 44

Correct Answer: C — Level 3 — Strategic and Crisis Issues

Explanation:

Level 3 covers **public harm potential, regulatory enforcement, major incidents, and reputational**

damage. The responsible stakeholders include C-suite leadership, board committees, and crisis-management teams.

Question 45

Correct Answer: B — Treat it as a Level 1 operational issue

Explanation:

Level 1 covers operational issues such as **performance degradation, minor bias detection, and routine maintenance.** The responsible teams include AI operations, system administrators, and operations analysts, with a response target of less than one hour in the study guide's illustrative model.

Question 46

Correct Answer: B — Level 2; AI governance committee, compliance officers, and legal counsel

Explanation:

Level 2 covers **policy violations, ethics concerns, compliance issues, and vendor problems.** The responsible teams include AI governance committees, compliance officers, and legal counsel.

Question 47

Correct Answer: B — Persistent issues and customer complaints can trigger escalation

Explanation:

Level 1 issues do not necessarily remain at Level 1 indefinitely. The study guide specifically identifies **persistent issues and customer complaints** as escalation triggers.

This illustrates an important governance principle: **severity can evolve during an incident**, so escalation criteria must be dynamic.

Question 48

Correct Answer: C — Strategic decisions, external communications, legal action, and business continuity

Explanation:

Level 3 provides authority for strategic and crisis-level decisions, including **external communications, legal action, and business continuity.** The escalation triggers include media attention, regulatory action, legal liability, and safety concerns.

Question 49

Correct Answer: B — Model purpose, intended applications, target users, and deployment context

Explanation:

The Model Overview and Intended Use section establishes the model's basic identity and purpose and provides stakeholders with context about how the model is intended to be used. This is distinct from risk-card information, which focuses on formal risk classification and impact assessment.

Question 50

Correct Answer: B — It provides structured information about the model, its intended use, limitations, and relevant characteristics

Explanation:

Model cards are part of the Module 6 **comprehensive AI documentation framework**. They support transparency, accountability, and informed evaluation. They do **not** replace risk assessments or continuous monitoring.

Question 51

Correct Answer: B — Data instance types, representativeness, missing information, demographic coverage, and geographic coverage

Explanation:

The **Composition Analysis** portion of dataset documentation examines data instance types and formats, representativeness, missing information, confidentiality/privacy considerations, and demographic and geographic coverage. These factors help identify quality and representation gaps.

Question 52

Correct Answer: A — Collection Process Documentation

Explanation:

Collection Process Documentation records **data acquisition methods, sampling strategies, temporal coverage, ethical review, quality assurance, and validation procedures**. These details help stakeholders understand how collection methods may have introduced bias or quality issues.

Question 53

Correct Answer: B — To support reproducibility and help downstream users understand potential effects on model behavior and performance

Explanation:

Preprocessing and cleaning documentation records transformation, normalization, software tools, algorithms, quality-control procedures, and validation methods. This supports **reproducibility** and helps downstream users understand how preprocessing can affect model performance and behavior.

Question 54

Correct Answer: A — Access procedures, licensing restrictions, intellectual-property considerations, and retention/deletion policies

Explanation:

Distribution and Access Controls documentation covers:

- Distribution mechanisms
 - Access procedures
 - Licensing terms
 - Usage restrictions
 - Copyright/IP considerations
 - Data retention and deletion policies
-

Question 55

Correct Answer: B — Likelihood and severity, health and safety risks, fundamental-rights impacts, societal/environmental effects, and residual risk

Explanation:

The **Risk Profile Assessment** section evaluates multiple impact dimensions. Module 6 specifically identifies likelihood/severity, health and safety, fundamental rights, societal/environmental effects, and residual risk after mitigation.

Question 56

Correct Answer: B — End-user interaction, human oversight requirements, consent procedures, feedback mechanisms, and complaint handling

Explanation:

The **Human Involvement and Control Mechanisms** section documents how humans interact with and control the AI system. It includes human oversight, subject involvement and consent, feedback mechanisms, and complaint handling.

Question 57

Correct Answer: B — Conduct comprehensive stakeholder engagement

Explanation:

The policy-development lifecycle begins with **comprehensive stakeholder engagement**. Module 6 recommends establishing cross-functional working groups so that legal, compliance, IT, business, HR, and potentially affected external stakeholders contribute to policy development.

Question 58

Correct Answer: B — Legal, compliance, IT, relevant business units, HR, and potentially affected external stakeholders

Explanation:

The module recommends broad cross-functional participation. This is important because AI governance crosses technical, legal, compliance, business, human-resource, and stakeholder-impact boundaries.

Question 59

Correct Answer: B — Use principle-based policies combined with specific requirements for critical risk and compliance areas

Explanation:

Module 6 recommends **principle-based policies** because they remain relevant longer as technology evolves. However, they should still establish specific requirements for critical risk and compliance areas. Practical implementation guidance and use-case scenarios should supplement the policy.

Question 60

Correct Answer: B — Policy and configuration changes, operational activity, incidents, exceptions, remediation, and regulatory inquiries

Explanation:

A comprehensive audit trail should capture multiple categories of governance activity, including:

- Policy modifications and approval workflows
- System configuration and parameter changes
- Access-control modifications
- Vendor/third-party changes
- AI usage and user interactions
- Model inference requests and responses
- Data access and processing
- Performance alerts
- Policy violations
- Remediation
- Incidents and response
- System failures and recovery
- Regulatory inquiries

The key concept is that an AI audit trail should provide **end-to-end evidence of governance and operational activity**, not merely application logs.

Question 61

An organization has implemented AI governance policies but reviews them only once every three years. New AI regulations and organizational AI use cases are emerging rapidly. Which approach BEST aligns with Module 6?

- A. Keep the policies unchanged to maintain consistency
 - B. Implement a continuous improvement loop involving feedback, changes, analysis, and performance monitoring
 - C. Replace all governance policies annually regardless of regulatory changes
 - D. Allow each development team to modify policies independently
-

Question 62

An AI governance platform generates alerts when model performance deteriorates and automatically routes the issue to the responsible team. Which capability does this MOST directly demonstrate?

- A. Continuous monitoring and automated governance
 - B. Data localization
 - C. Regulatory arbitrage
 - D. Model documentation
-

Question 63

An organization wants to detect whether an AI model's accuracy is declining because the underlying data distribution has changed. Which capability should it prioritize?

- A. Data retention monitoring
 - B. Model drift detection
 - C. Contract management
 - D. Policy version control only
-

Question 64

An AI system produces materially different outcomes for two demographic groups despite having acceptable overall accuracy. Which monitoring capability is MOST important?

- A. Bias detection across relevant populations
 - B. Cloud cost monitoring
 - C. Software license monitoring
 - D. Vendor invoice reconciliation
-

Question 65

An organization wants its audit trail to demonstrate that an AI governance decision was properly approved. Which information would be MOST important to capture?

- A. Only the final decision
 - B. Policy modification, approval workflow, relevant configuration changes, and associated governance activity
 - C. Only the model's inference results
 - D. Only the names of system administrators
-

Question 66

A governance team wants to investigate why a production AI system changed its behavior after a configuration update. Which audit-trail capability would be MOST valuable?

- A. Recording system configuration changes and parameter updates
 - B. Recording only annual compliance assessments
 - C. Recording only employee training attendance
 - D. Recording only customer complaints
-

Question 67

A company operates an AI platform in Europe but processes training and inference data through cloud infrastructure located in multiple countries. Which issue should the organization assess FIRST?

- A. Whether the AI system has enough computing power
 - B. Cross-border data flows and the legal mechanisms supporting those transfers
 - C. Whether all developers use the same programming language
 - D. Whether the AI model is open source
-

Question 68

A multinational organization uses Standard Contractual Clauses (SCCs) for an international AI data transfer. According to Module 6, what additional assessment is associated with SCC-based transfers?

- A. Transfer Impact Assessment
 - B. Model Accuracy Assessment
 - C. AI Safety Officer Assessment
 - D. Software License Assessment
-

Question 69

A multinational corporation has complex intra-group AI data transfers between subsidiaries in several jurisdictions. Which mechanism described in Module 6 may provide an appropriate framework for these intra-group transfers?

- A. Binding Corporate Rules
- B. Model Cards
- C. MITRE ATLAS
- D. AI-CAIQ

Question 70

An organization wants to reduce the need to centralize sensitive data across national borders while still training an AI model. Which technical approach described in Module 6 is MOST appropriate?

- A. Federated learning
 - B. Centralized data replication
 - C. Permanent data retention
 - D. Regulatory arbitrage
-

Question 71

A company chooses to conduct AI development in a country solely because its AI data regulations are less restrictive, allowing the company to use data that would not be permitted in another market where the AI service will eventually operate. What governance challenge does this illustrate?

- A. Data lineage
 - B. Regulatory arbitrage
 - C. Model drift
 - D. Conformity assessment
-

Question 72

A U.S.-based company primarily serves U.S. customers but must redesign aspects of its AI system because it affects individuals in the EU. Which global AI governance challenge does this illustrate?

- A. Data retention
 - B. Extraterritorial application
 - C. Model card incompleteness
 - D. Vendor concentration risk
-

Question 73

A country requires certain categories of data to remain within its national borders. A multinational organization wants to train a global AI model using that data. Which challenge is MOST directly involved?

- A. Data localization
 - B. Model drift
 - C. RACI ambiguity
 - D. AI committee composition
-

Question 74

An organization wants to transfer a dual-use AI technology internationally. The technology has both commercial and potential military applications. Which issue should receive particular attention?

- A. Export controls and national security considerations
 - B. Model-card formatting
 - C. Employee training attendance
 - D. AI governance committee size
-

Question 75

An organization is beginning an AI governance program and wants an immediate compliance-oriented action plan before building its long-term framework. According to the Module 6 30/60/90-day framework, what should be targeted during the **first 30 days**?

- A. Advanced predictive risk analytics
 - B. Immediate compliance
 - C. Full enterprise implementation
 - D. Certification against ISO 42001
-

Question 76

Under the Module 6 30/60/90-day framework, what is the principal objective for approximately **60 days**?

- A. Establish a policy and framework
 - B. Replace all existing AI systems
 - C. Complete international certification
 - D. Deploy predictive analytics
-

Question 77

A company has completed its initial 30-day compliance assessment and established its AI governance framework during the next phase. What is the stated objective for approximately **90 days**?

- A. Begin regulatory research
 - B. Full implementation
 - C. Identify the need for AI governance
 - D. Conduct only a preliminary stakeholder meeting
-

Question 78

An organization has entered the first two months of its formal 12-month AI governance roadmap. Which activity BEST fits this phase?

- A. Assessment and framework selection based on organizational needs
- B. Advanced AI-driven predictive risk analytics
- C. Benchmarking governance maturity against industry peers
- D. Pursuing external certification as the primary activity

Question 79

An organization is in **months 7–9** of the Module 6 implementation roadmap. Which activity BEST fits this phase?

- A. Initial framework selection
 - B. Core framework implementation
 - C. Integration with existing operational processes and tools
 - D. Only regulatory research
-

Question 80

A large organization has more than 10,000 employees and is planning a comprehensive AI governance implementation. Which staffing and timeline combination BEST matches the Module 6 guidance?

- A. 2–3 FTE and 6–9 months
 - B. 5–7 FTE and 9–12 months
 - C. 10–15 FTE and 12–18 months
 - D. 20–30 FTE and 24–36 months
-

Answers & Explanations**Question 61**

Correct Answer: B — Implement a continuous improvement loop involving feedback, changes, analysis, and performance monitoring

Explanation:

Module 6 describes a continuous improvement loop consisting of:

1. Collect feedback
2. Implement changes based on feedback and new regulations
3. Analyze performance
4. Monitor performance at an overarching level

This is particularly important because AI technology and governance requirements can change rapidly.

Question 62

Correct Answer: A — Continuous monitoring and automated governance

Explanation:

Module 6 describes advanced governance capabilities that include **automated governance tools for**

continuous compliance monitoring. Automated workflows and monitoring can detect issues and initiate appropriate governance processes without relying exclusively on manual review.

Question 63

Correct Answer: B — Model drift detection

Explanation:

Model drift refers to a decline in AI model performance over time caused by changes in data patterns. Module 6 identifies **drift detection capabilities** as part of quality and performance monitoring and emphasizes continuous monitoring and potential retraining.

Question 64

Correct Answer: A — Bias detection across relevant populations

Explanation:

Module 6 identifies continuous monitoring as a mechanism for detecting **drift and bias**. Risk documentation also requires accuracy assessment across different conditions and populations, together with robustness and performance monitoring.

The important distinction is that **overall accuracy alone may not demonstrate equitable performance across populations**.

Question 65

Correct Answer: B — Policy modification, approval workflow, relevant configuration changes, and associated governance activity

Explanation:

A comprehensive audit trail must capture governance activity, not simply the final outcome. Module 6 specifically includes **governance policy modifications with approval workflows**, system configuration changes, access-control modifications, vendor changes, operational activity, incidents, and compliance activity.

Question 66

Correct Answer: A — Recording system configuration changes and parameter updates

Explanation:

Configuration and parameter changes are explicitly included in the comprehensive audit-trail requirements. This evidence can help reconstruct what changed and correlate that change with subsequent system behavior.

Question 67

Correct Answer: B — Cross-border data flows and the legal mechanisms supporting those transfers

Explanation:

Cross-border AI operations require organizations to understand where data is processed and what legal mechanisms support international transfers. Module 6 emphasizes mapping cross-border training, inference, and model-update data flows and conducting transfer impact assessments.

Question 68

Correct Answer: A — Transfer Impact Assessment

Explanation:

Module 6 states that **Standard Contractual Clauses (SCCs)** are commonly used for transfers to non-adequate jurisdictions and require **Transfer Impact Assessments (TIAs)** following Schrems II requirements. Additional safeguards may be necessary depending on the destination country's surveillance laws.

Question 69

Correct Answer: A — Binding Corporate Rules

Explanation:

Binding Corporate Rules (BCRs) are described as internal data-protection rules for multinational organizations with complex international data flows. They provide a framework for **intra-group transfers** and require approval from the relevant supervisory authority.

Question 70

Correct Answer: A — Federated learning

Explanation:

Federated learning allows models to be trained without centralizing sensitive data across borders. Module 6 identifies it as one of the privacy-preserving techniques that can help address cross-border AI data challenges, alongside differential privacy and synthetic data.

Question 71

Correct Answer: B — Regulatory arbitrage

Explanation:

Regulatory arbitrage occurs when organizations seek jurisdictions or regulatory environments with lighter requirements in order to conduct activities that may be restricted elsewhere. Module 6 identifies this as a major global AI governance challenge.

Question 72

Correct Answer: B — Extraterritorial application

Explanation:

Extraterritorial application occurs when an organization's activities outside a jurisdiction are nevertheless affected by that jurisdiction's laws because its AI systems impact individuals or markets there. Module 6 uses the example of U.S. companies needing to modify AI systems to comply with EU requirements.

Question 73

Correct Answer: A — Data localization

Explanation:

Data localization requirements require certain categories of data to remain within particular jurisdictions. Module 6 identifies these requirements as a significant challenge for organizations attempting to develop and deploy AI globally.

Question 74

Correct Answer: A — Export controls and national security considerations

Explanation:

Module 6 identifies **export controls** as a significant challenge for dual-use AI technologies because such technologies may have both commercial and military applications. National-security considerations can therefore affect international transfers.

Question 75

Correct Answer: B — Immediate compliance

Explanation:

The Module 6 30/60/90-day framework is:

- **30 days → Immediate compliance**
- **60 days → Establish a policy and framework**
- **90 days → Full implementation**

This framework is intended to provide immediate action while the organization develops a more comprehensive governance program.

Question 76

Correct Answer: A — Establish a policy and framework

Explanation:

The second milestone in the 30/60/90-day model is establishing the organization's **policy and governance framework**.

Question 77

Correct Answer: B — Full implementation

Explanation:

The third milestone is **full implementation at approximately 90 days**. This should be understood as the stated immediate-action framework; the study guide separately provides a more comprehensive 12-month implementation strategy for mature enterprise deployment.

Question 78

Correct Answer: A — Assessment and framework selection based on organizational needs

Explanation:

The 12-month roadmap identifies the first phase as:

Months 1–2: Assessment and framework selection based on organizational needs.

The organization should assess its needs and determine the appropriate governance framework before moving into core implementation.

Question 79

Correct Answer: C — Integration with existing operational processes and tools

Explanation:

The roadmap specifies:

- **Months 1–2:** Assessment and framework selection
- **Months 3–6:** Core framework implementation
- **Months 7–9:** Integration with existing operational processes and tools
- **Months 10–12:** Advanced capabilities and continuous improvement

Therefore, months 7–9 focus on **integration** rather than initial framework selection or advanced optimization.

Question 80

Correct Answer: C — 10–15 FTE and 12–18 months

Explanation:

Module 6 provides different implementation assumptions based on organizational size:

Organization	Timeline	Staffing
Small (<1,000)	6–9 months	2–3 FTE
Medium (1,000–10,000)	9–12 months	5–7 FTE

Organization	Timeline	Staffing
Large (>10,000)	12–18 months	10–15 FTE

Question 81

An organization wants an accountability framework that emphasizes four complementary areas: clear organizational roles, data quality, continuous performance assessment, and real-time oversight. Which framework described in Module 6 BEST matches this requirement?

- A. MITRE ATLAS
 - B. GAO AI Accountability Framework
 - C. ISO/IEC 23894
 - D. CSA AICM
-

Question 82

A company's AI governance team is reviewing the four principles of the GAO AI Accountability Framework. Which activity BEST represents the **Data** principle?

- A. Establishing clear executive accountability for AI decisions
 - B. Assessing data quality and integrity throughout the AI lifecycle
 - C. Establishing real-time incident monitoring dashboards
 - D. Defining the organization's AI risk tolerance
-

Question 83

A company wants its AI governance framework to focus specifically on **threat intelligence, adversarial tactics, and security analysis of AI systems**. Which combination of framework and purpose is MOST appropriate?

- A. MITRE ATLAS — AI threat landscape and security
 - B. ISO 42001 — AI management-system certification
 - C. ISO 23894 — AI risk management methodology
 - D. NIST AI RMF — cloud infrastructure controls
-

Question 84

A multinational organization wants to establish a **formally certifiable AI management system** that can support its global AI governance program. Which framework should it prioritize?

- A. NIST AI RMF
- B. ISO/IEC 23894
- C. ISO/IEC 42001
- D. MITRE ATLAS

Question 85

An organization wants internationally recognized AI risk-management guidance but does not currently require certification. Which framework is the BEST fit?

- A. ISO/IEC 23894
 - B. ISO/IEC 42001
 - C. STAR for AI Level 2
 - D. MITRE ATLAS
-

Question 86

A cloud service provider is primarily concerned with implementing structured security controls for its cloud-based AI services and potentially demonstrating assurance through CSA's STAR for AI program. Which combination BEST fits this requirement?

- A. CSA AICM and STAR for AI
 - B. MITRE ATLAS and ISO 23894
 - C. NIST AI RMF and GDPR
 - D. ISO 42001 and California ADMT
-

Question 87

An organization is reviewing its RACI matrix for AI governance. A project has three people marked as "Accountable." What is the PRIMARY governance concern?

- A. Too many people are marked Responsible
 - B. Accountability for final decision-making is ambiguous
 - C. Too many people are marked Informed
 - D. The project has insufficient technical expertise
-

Question 88

A governance manager wants to ensure that the RACI matrix remains effective over time. Which combination BEST reflects the Module 6 guidance?

- A. Rotate the Accountable person regularly and maximize the number of Consulted stakeholders
 - B. Maintain a single point of accountability, integrate with existing frameworks, and prevent individual overload
 - C. Assign everyone as Responsible to ensure maximum ownership
 - D. Remove all Informed stakeholders to simplify communication
-

Question 89

An organization has limited AI governance resources and relatively limited AI adoption. Leadership is considering assigning AI governance responsibilities to an existing technology executive who already serves as both CTO and CDO. What concern should leadership consider according to Module 6?

- A. The structure is prohibited under AI governance principles
 - B. The dual-hat approach can be cost-effective but may create competing priorities and limited governance focus
 - C. The AI Safety Officer must always report directly to the board
 - D. AI governance cannot be assigned to existing executives
-

Question 90

An organization considers AI a core strategic capability and recognizes significant enterprise-level AI risk. Which AI Safety Officer reporting structure provides the **maximum organizational authority and resource access** according to Module 6?

- A. Dual-hat CTO/CDO role
 - B. CRO report
 - C. Direct CEO/COO report
 - D. Business-unit manager report
-

Question 91

An organization wants its AI governance program to operate effectively as a **second line of defense**. Which combination of functions should the AI oversight role leverage?

- A. Only software engineering and product development
 - B. Compliance, risk management, cybersecurity, and privacy functions
 - C. Only external auditors
 - D. Only business-unit leadership
-

Question 92

A company's AI governance committee wants to determine whether its program is actually improving rather than merely producing documentation. Which measurement approach BEST aligns with Module 6?

- A. Measure only the number of policies created
 - B. Track documentation, training, governance, risk management, compliance, and technology dimensions
 - C. Measure only AI model accuracy
 - D. Measure only the number of AI systems deployed
-

Question 93

An AI governance dashboard shows that AI inventory accuracy is 98%, but the percentage of AI systems with completed risk assessments is only 40%. Which category of accountability metrics does the incomplete risk-assessment rate represent?

- A. Coverage metrics
 - B. Compliance metrics
 - C. Performance metrics
 - D. Financial metrics
-

Question 94

An AI governance committee receives operational metrics continuously, tactical summaries weekly, strategic summaries monthly, and comprehensive assessments quarterly. Which statement BEST describes this reporting structure?

- A. It uses different reporting cadences for different levels of governance and decision-making
 - B. It indicates inconsistent governance reporting
 - C. It should provide all stakeholders with the same report once a year
 - D. It eliminates the need for annual program assessment
-

Question 95

A company is integrating an external generative AI provider into a customer-facing application. Which governance action BEST reflects the Module 6 implementation guidance?

- A. Treat the provider as outside the organization's governance scope
 - B. Enhance third-party risk management for AI services and integrate vendor controls into the governance program
 - C. Rely exclusively on the vendor's marketing claims
 - D. Eliminate internal monitoring because the vendor owns the model
-

Question 96

An organization is in the **framework implementation and integration** stage of its AI governance roadmap. Which activity BEST belongs in this phase?

- A. Conduct only the initial AI inventory
 - B. Deploy RACI matrices, integrate governance with existing risk programs, implement monitoring, and enhance vendor management
 - C. Focus exclusively on future AI technologies
 - D. Perform only industry benchmarking
-

Question 97

An organization has reached **months 10–12** of its implementation roadmap. Which activity BEST reflects the focus of this phase?

- A. Initial framework selection
 - B. Basic AI inventory development
 - C. Predictive risk management, advanced governance analytics, and continuous improvement
 - D. Initial stakeholder identification only
-

Question 98

A large organization has strong executive sponsorship, but its AI governance program remains ineffective because IT, legal, compliance, and business units operate independently. Which critical success factor is MOST clearly lacking?

- A. Cross-functional collaboration
 - B. Technology investment
 - C. Certification
 - D. Model-card adoption
-

Question 99

An organization wants to demonstrate that its AI governance program is improving its ability to identify and reduce AI risk. Which combination of metrics would provide the BEST evidence?

- A. Number of AI models purchased and annual cloud expenditure
 - B. Risk-reduction trends, control effectiveness, regulatory compliance, incident response time, and resolution effectiveness
 - C. Number of developers hired and amount of training data
 - D. Number of AI experiments conducted and model inference speed
-

Question 100

A multinational financial institution is implementing an enterprise AI governance program. It inventories AI systems, classifies risks, establishes an AI Safety Officer, creates RACI matrices, implements model/data/risk documentation, integrates monitoring with existing infrastructure, manages cross-border data issues, and establishes continuous improvement. Which statement BEST describes this approach?

- A. It is primarily an AI model-development program
 - B. It represents an integrated AI governance, risk management, and compliance program across the AI lifecycle
 - C. It is primarily a cybersecurity incident-response program
 - D. It is primarily a data-science quality-assurance program
-

Answers & Explanations

Question 81

Correct Answer: B — GAO AI Accountability Framework

Explanation:

The Module 6 vocabulary identifies four complementary GAO AI Accountability principles:

- **Governance** — clear roles and accountability
- **Data** — quality and integrity
- **Performance** — continuous monitoring
- **Monitoring** — real-time oversight and feedback loops

The key exam clue is the combination of **governance + data + performance + monitoring**.

Question 82

Correct Answer: B — Assessing data quality and integrity throughout the AI lifecycle

Explanation:

The **Data** principle of the GAO framework focuses on data quality and integrity. It complements Governance, Performance, and Monitoring rather than replacing them.

Question 83

Correct Answer: A — MITRE ATLAS — AI threat landscape and security

Explanation:

MITRE ATLAS is specifically positioned around the **AI adversarial threat landscape and security**. The Module 6 vocabulary describes ATLAS as extending MITRE ATT&CK to AI systems, with 56 techniques across 14 adversarial tactics.

A useful distinction:

- **ATLAS** → threats/adversarial security
 - **NIST AI RMF** → comprehensive AI risk management
 - **ISO 23894** → AI risk management
 - **ISO 42001** → AI management system
-

Question 84

Correct Answer: C — ISO/IEC 42001

Explanation:

The framework selection matrix identifies **ISO 42001** as an AI management-system standard with certification available. It is particularly suited to organizations requiring a formal, certifiable AI management system.

Question 85

Correct Answer: A — ISO/IEC 23894

Explanation:

ISO/IEC 23894 provides an internationally recognized **AI risk-management methodology**. The study guide specifically distinguishes it from ISO 42001: ISO 23894 provides risk-management guidance, while ISO 42001 provides certifiable AI management-system requirements.

Question 86

Correct Answer: A — CSA AICM and STAR for AI

Explanation:

The CSA AICM addresses **cloud and AI security controls**, while STAR for AI operationalizes those controls through an assurance model. The study guide's recommended progression is:

AI-CAIQ → AICM controls → STAR certification/assurance.

Question 87

Correct Answer: B — Accountability for final decision-making is ambiguous

Explanation:

In RACI:

- Multiple people can be **Responsible**
- Only **one person should be Accountable** for a task
- Consulted people provide input
- Informed people receive updates

The purpose of a single Accountable party is to avoid ambiguity over final authority.

Question 88

Correct Answer: B — Maintain a single point of accountability, integrate with existing frameworks, and prevent individual overload

Explanation:

The sample questions supplied with the study guide emphasize exactly this principle. RACI reviews should preserve **single-point accountability**, integrate with existing governance structures, and maintain balanced workload distribution rather than creating accountability overload.

Question 89

Correct Answer: B — The dual-hat approach can be cost-effective but may create competing priorities and limited governance focus

Explanation:

The Module 6 organizational-placement section describes the **dual-hat CTO/CDO role** as common where AI governance resources are limited. Its advantage is cost-effectiveness, but the potential disadvantages are **competing priorities and limited governance focus**.

Question 90

Correct Answer: C — Direct CEO/COO report

Explanation:

For organizations where AI is a **core strategic capability or represents significant enterprise risk**, Module 6 identifies the direct CEO/COO reporting structure as appropriate because it provides **maximum organizational authority and resource access**.

Question 91

Correct Answer: B — Compliance, risk management, cybersecurity, and privacy functions

Explanation:

The study guide positions AI oversight within the **second line of defense** and recommends leveraging existing compliance, risk management, cybersecurity, and privacy functions while addressing AI-specific requirements that traditional oversight may not adequately cover.

Question 92

Correct Answer: B — Track documentation, training, governance, risk management, compliance, and technology dimensions

Explanation:

Module 6 describes a radar-chart approach for measuring governance success across six categories:

1. Documentation
2. Training
3. Governance
4. Risk management
5. Compliance
6. Technology

The important point is that AI governance maturity is **multidimensional**.

Question 93

Correct Answer: A — Coverage metrics

Explanation:

The Module 6 accountability metrics identify **AI system inventory accuracy and risk-assessment completion rates** as coverage metrics.

Therefore:

- Inventory accuracy → Coverage
 - Risk-assessment completion → Coverage
 - Policy adherence/training/audit findings → Compliance
 - Incident response/resolution/customer satisfaction → Performance
 - Risk reduction/control effectiveness/regulatory compliance → Risk
-

Question 94

Correct Answer: A — It uses different reporting cadences for different levels of governance and decision-making

Explanation:

Module 6 provides the following reporting cadence:

Frequency Primary audience

Real-time Governance committee and operations

Weekly Governance committee and technical leadership

Monthly Executive leadership and risk committees

Quarterly Board committees and external stakeholders

Annual Program effectiveness assessment

This creates **tiered reporting**, allowing operational issues and strategic governance matters to be handled at appropriate levels.

Question 95

Correct Answer: B — Enhance third-party risk management for AI services and integrate vendor controls into the governance program

Explanation:

Module 6 explicitly identifies **enhanced third-party risk management for AI services** as part of the framework implementation phase. AI vendors remain an important component of the organization's overall AI risk landscape.

The governance principle is that outsourcing an AI capability does **not automatically outsource the organization's governance responsibility**.

Question 96

Correct Answer: B — Deploy RACI matrices, integrate governance with existing risk programs, implement monitoring, and enhance vendor management

Explanation:

The framework implementation and integration phase includes:

- AI governance policy development
 - RACI implementation
 - Integration with existing risk-management/compliance programs
 - Training
 - Governance-platform integration
 - Audit-trail and compliance monitoring
 - Automated compliance workflows
 - AI-specific incident response
 - Enhanced AI vendor management
 - Model cards, data sheets, and risk cards
 - Compliance verification and stakeholder reporting
-

Question 97

Correct Answer: C — Predictive risk management, advanced governance analytics, and continuous improvement

Explanation:

Months **10–12** represent the optimization and continuous-improvement phase. The study guide identifies:

- AI-driven governance
- Continuous compliance monitoring
- Predictive risk management
- Early-warning systems
- Advanced governance analytics
- Integration optimization
- External stakeholder engagement
- Third-party assessments
- Certification evaluation
- Benchmarking

- Continuous improvement
-

Question 98

Correct Answer: A — Cross-functional collaboration

Explanation:

The study guide identifies **cross-functional collaboration** as a critical success factor, specifically requiring effective integration across IT, legal, compliance, and business functions.

Strong executive sponsorship alone cannot compensate for organizational silos.

Question 99

Correct Answer: B — Risk-reduction trends, control effectiveness, regulatory compliance, incident response time, and resolution effectiveness

Explanation:

These measures directly correspond to Module 6's **risk and performance metrics**:

- Risk-reduction trends
- Control effectiveness
- Regulatory compliance
- Incident response times
- Resolution effectiveness
- Customer satisfaction

This is much stronger evidence of governance effectiveness than simply counting AI models, developers, or infrastructure expenditure.

Question 100

Correct Answer: B — It represents an integrated AI governance, risk management, and compliance program across the AI lifecycle

Explanation:

This scenario combines essentially every major capability developed throughout Module 6:

- AI inventory and risk classification
- AI Safety Officer and organizational accountability
- RACI
- Model cards, data sheets, and risk cards
- Continuous monitoring

- Audit trails
- Existing infrastructure integration
- Cross-border governance
- Risk management
- Compliance
- Continuous improvement

The Module 6 summary explicitly describes AI governance as going beyond traditional IT governance to address **fairness, transparency, accountability, societal impact, risk management, compliance, documentation, and continuous monitoring**